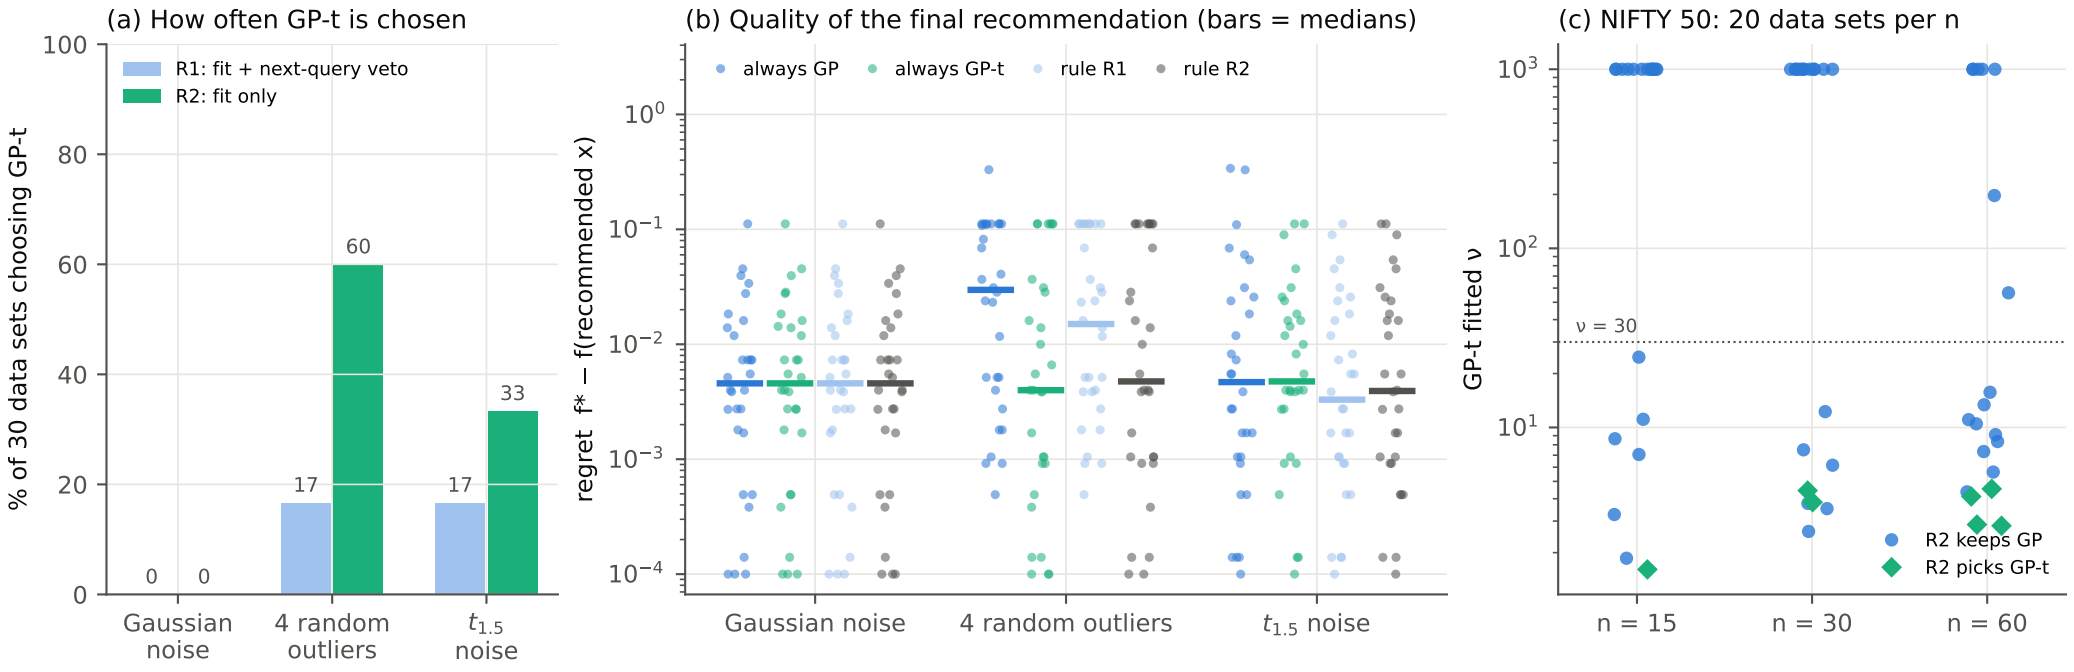
(a) How often GP-t is chosen
% of 30 data sets choosing GP-t
R1: fit + next-query veto
R2: fit only
0
0
17
60
17
33
Gaussian noise
4 random outliers
$t_{1.5}$ noise
(b) Quality of the final recommendation (bars = medians)
regret f* − f(recommended x)
always GP
always GP-t
rule R1
rule R2
Gaussian noise
4 random outliers
$t_{1.5}$ noise
(c) NIFTY 50: 20 data sets per n
GP-t fitted ν
ν = 30
R2 keeps GP
R2 picks GP-t
n = 15
n = 30
n = 60